

Non Parametric Filters

Histogram and Particle Filters


Department of Electrical and Electronics Engineering
Dr. Afşar Saranlı

*Lecture slides heavily use material from the textbook and
Sebastian Thrun, Lecture Slides; <http://www.probablistic-robotics.org/>*



What we will discuss



- Discuss the merits of non-parametric density representations,
- Remember histograms and introduce the *Histogram Filters*,
- Introduce *Particle Filters*,
- Introduce *resampling* and *Importance Sampling*,
- Review practical issues and properties of particle filters,
- Try to conclude...





Non-Parametric Representations



- What is a non-parametric density representation?
- **Finite set of values (samples) instead of a parametric closed form expression**
- Example: *Gaussian* versus its *histogram*





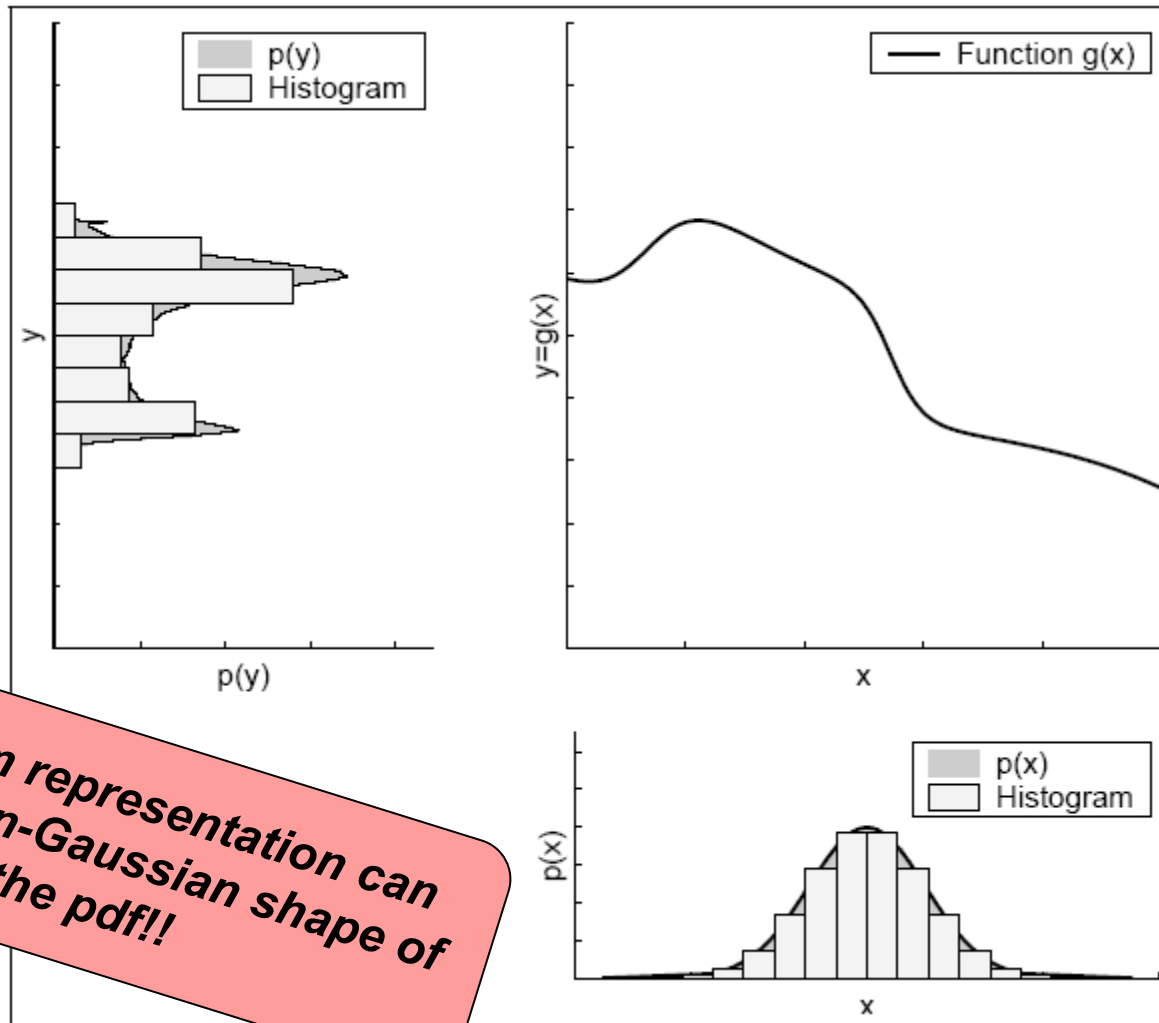
Why Non-Parametric?

- Why would we want to use non-parametric?
- Why would we not want to use non-parametric?
- Quality? Computational Complexity?
- As $N \rightarrow \infty$, non-parametric rep. converges uniformly to the true density.





What do you see?

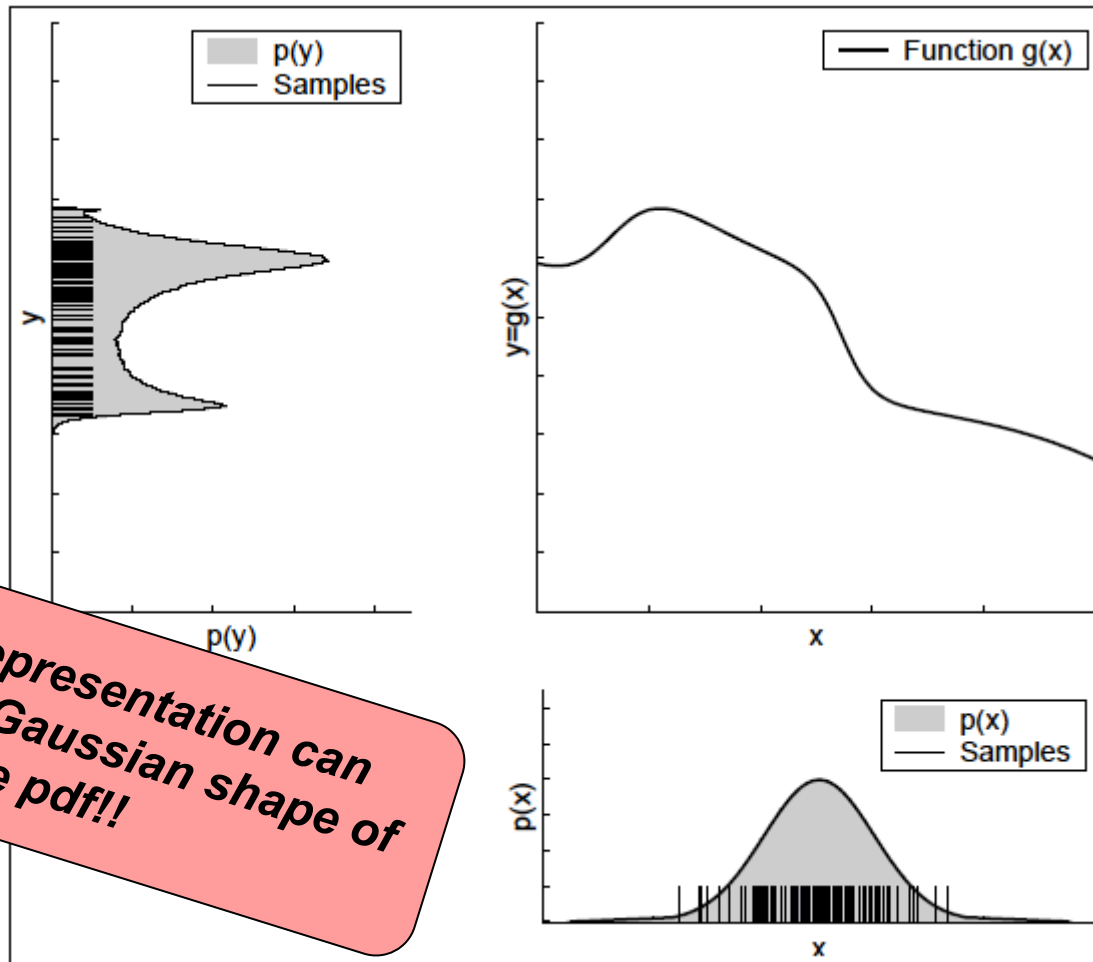


Histogram representation can capture Non-Gaussian shape of the pdf!!



Other Non-Parametric Reps

- **Particles (Sample based) Representation**



“Partile” representation can capture Non-Gaussian shape of the pdf!!



Non-Parametric Advantages

- Both histograms and particle sets have:
- No strong parametric assumptions on density (any arbitrary shape can be represented),
- Accuracy of the representation can be adjusted as required (by setting N)
- Results in conceptually much simpler program implementations,
- Well suited for *complex multi-modal beliefs* (e.g. in *global localization* with hard data association problems)



Non-Parametric Disadvantages

- **Simply: *Computational Complexity!!***
- A naïve implementation can be orders of magnitude more complex than parametric implementations (e.g. Kalman Filters)
- Problem becomes compounded if state-space dimension increases. (much larger N needed!)
- Fortunately:
 - Computational complexity can be adapted by adapting number of parameters N ,
 - Complexity and accuracy can be traded off using: *Resource Adaptation*,
 - Resource adaptive algorithms very important in robotics and embedded systems



Question

- How would we use non-parametric density representations to implement recursive state estimation?

***Use the Bayes Filtering framework
in discrete-form***



Part1: The Histogram Filter



Part1: The Histogram Filter

- Remember the generic Bayes Filter:

```
1:   Algorithm Bayes_filter( $bel(x_{t-1}), u_t, z_t$ ):  
2:     for all  $x_t$  do  
3:        $\overline{bel}(x_t) = \int p(x_t | u_t, x_{t-1}) bel(x_{t-1}) dx_{t-1}$   
4:        $bel(x_t) = \eta p(z_t | x_t) \overline{bel}(x_t)$   
5:     endfor  
6:     return  $bel(x_t)$ 
```

Table 2.1 The general algorithm for Bayes filtering.



The Histogram Filter - DBF

- Discrete Bayes Filter from the Bayes Filter

```
1:   Algorithm Bayes_filter( $bel(x_{t-1}), u_t, z_t$ ):  
2:     for all  $x_t$  do  
3:        $\overline{bel}(x_t) = \int p(x_t | u_t, x_{t-1}) bel(x_{t-1}) dx_{t-1}$   
4:        $bel(x_t) = \eta p(z_t | x_t) \overline{bel}(x_t)$   
5:     endfor  
6:     return  $bel(x_t)$ 
```

```
1:   Algorithm Discrete_Bayes_filter( $\{p_{k,t-1}\}, u_t, z_t$ ):  
2:     for all  $k$  do  
3:        $\bar{p}_{k,t} = \sum_i p(X_t = x_k | u_t, X_{t-1} = x_i) p_{i,t-1}$   
4:        $p_{k,t} = \eta p(z_t | X_t = x_k) \bar{p}_{k,t}$   
5:     endfor  
6:     return  $\{p_{k,t}\}$ 
```




The Histogram Filter - Discretization

- Some problems are already discrete in nature (state of a door, states of a game board),
- Some problems are represented in discrete form because of the *resolution of interest*. (E.g. occupancy grid maps)
- Entirely continuous problems may also be discretized through various approximations. (e.g. the orientation of a robot with 5° steps),
- Granularity (resolution) of the discretization may be very important (not only for performance but also for proper operation of the filters)



Example – Estimation of Door State

- **Problem is to estimate a single number!!**
 $\text{Prob}(\text{door}=\text{open}|\text{all past states, all past measurements})$





Continuous State Space

- When DBF is applied to a continuous state space, it is called the *Histogram Filter*,
- Histogram Filter decompose such a space into “bins” through a suitable “partitioning”,

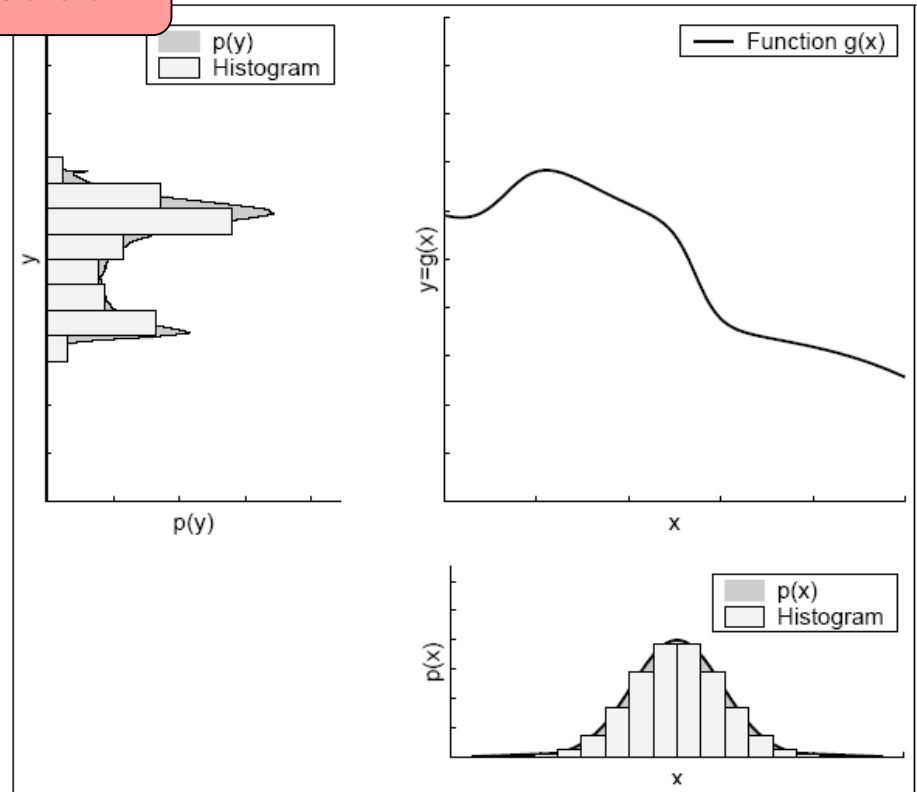
$$\begin{aligned}\text{dom}(X_t) &= \mathbf{x}_{1,t} \cup \mathbf{x}_{2,t} \cup \dots \mathbf{x}_{K,t} \\ \mathbf{x}_{i,t} \cap \mathbf{x}_{k,t} &= \emptyset \text{ and } \bigcup_k \mathbf{x}_{k,t} = \text{dom}(X_t)\end{aligned}$$

- Most common partitioning is a multi-dimensional grid representation,
- Resulting Pdf approximation is a piecewise constant pdf.




Example – 1D / 2D spaces

1D space



2D space





Use in Continuous Spaces

- We might be given the continuous densities,
 $p(x_t \mid u_t, x_{t-1})$ and $p(z_t \mid x_t)$
- These are defined for a continuum of states, not for the discrete “bins”,
- How can we discretize the given continuous densities?



Use in Continuous Spaces

- Given continuous $p(x_t \mid u_t, x_{t-1})$ and $p(z_t \mid x_t)$
- For each bin, we can pick a *representative “mean” state*:

$$\hat{x}_{k,t} = |\mathbf{x}_{k,t}|^{-1} \int_{\mathbf{x}_{k,t}} x_t dx_t$$

- *Then approximate the discrete probability mass functions as:*

$$\begin{aligned} p(z_t \mid \mathbf{x}_{k,t}) &\approx p(z_t \mid \hat{x}_{k,t}) \\ p(\mathbf{x}_{k,t} \mid u_t, \mathbf{x}_{i,t-1}) &\approx \eta |\mathbf{x}_{k,t}| p(\hat{x}_{k,t} \mid u_t, \hat{x}_{i,t-1}) \end{aligned}$$

There is a correction in the book here.

Discrete “probability” is calculated from the continuous “likelihood” value at the mean state which is integrated over the “bin”



Use in Continuous Spaces

- Then, the Discrete Bayes Filter can be used directly:

```
1:   Algorithm Discrete_Bayes_filter( $\{p_{k,t-1}\}, u_t, z_t$ ):  
2:     for all  $k$  do  
3:        $\bar{p}_{k,t} = \sum_i p(X_t = x_k \mid u_t, X_{t-1} = x_i) p_{i,t-1}$   
4:        $p_{k,t} = \eta p(z_t \mid X_t = x_k) \bar{p}_{k,t}$   
5:     endfor  
6:     return  $\{p_{k,t}\}$ 
```

$$p(z_t \mid \mathbf{x}_{k,t}) \approx p(z_t \mid \hat{x}_{k,t})$$
$$p(\mathbf{x}_{k,t} \mid u_t, \mathbf{x}_{i,t-1}) \approx \eta |\mathbf{x}_{k,t}| p(\hat{x}_{k,t} \mid u_t, \hat{x}_{i,t-1})$$



Practical Issues: Decomposition

- Histogram filters can trade-off *accuracy* with *computational complexity* but...
- Desired accuracy may come at a prohibitive computational price!!
- A Naïve uniform grid decomposition with full update may be unusable.
- Some ideas:
 - Density trees,
 - Selective Updating,
 - “Topological” representations,


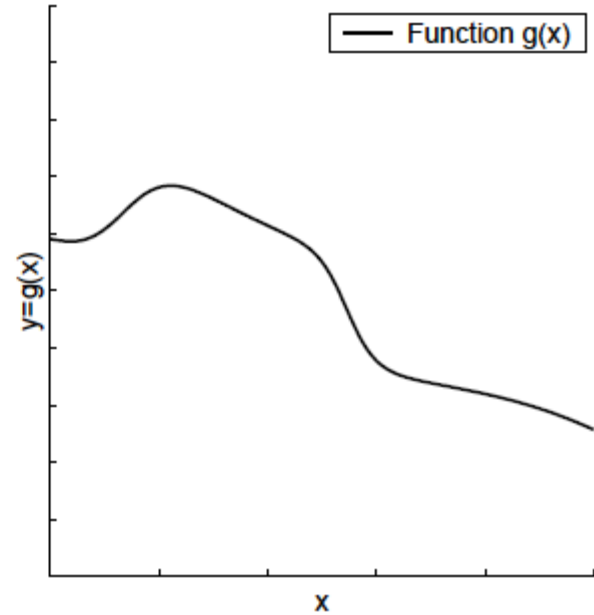




Decomposition: Density Trees

- Important example of “dynamic decomposition”,
- Dynamic techniques adapt to the shape of the posterior being approximated,
- Static techniques *easier to implement*, Dynamic techniques *more efficient and hence faster*,
- ***Density Trees***: A recursive decomposition that takes into account the distribution,
 - *The more likely a region, the finer the decomposition (more bins) and vice versa.*
 - *Achieves higher approximation quality with the same computational complexity,*
 - *OR: Cuts the complexity by orders of magnitude for the same approximation quality*



Dynamic Decomposition: Example






Practical Issues: Selective Update

- Instead of worrying about the decomposition, worry about the filter updates,
- E.g., generate a uniform grid decomposition but...
- Update only a fraction of the cells (probabilities) at a time, specifically...
- **Those that have the posterior probabilities that exceeds a certain threshold.**
- It can also be viewed as a dynamic decomposition technique,
- Can save orders of magnitude in computational complexity. What about storage space?



Practical Issues: Topological Reps

- “Metric” vs “Topological” representations,
- **Topological:** Coarse, graph-like representations where only *significant places or features* are stored,
- E.g. corridors, intersections, dead ends,
- “Topological” representations usually more efficient but much less precise,





Special: Binary Bayes w Static State

- A special case of Discrete Bayes Filter,
- Best approximation for certain problems (e.g. occupancy grid maps)
- State is static: belief at time t is only a function of the measurements:

$$bel_t(x) = p(x \mid z_{1:t}, u_{1:t}) = p(x \mid z_{1:t})$$

- An elegant and efficient formulation using the so called “log-odds ratios”
- Uses the inverse measurement model $p(x|z_t)$ contrary to our usual model $p(z_t|x)$



Special: Binary Bayes w Static State

$P(x)$: Prior probability of state x

```
1: Algorithm binary_Bayes_filter( $l_{t-1}, z_t$ ):  
2:    $l_t = l_{t-1} + \log \frac{p(x|z_t)}{1-p(x|z_t)} - \log \frac{p(x)}{1-p(x)}$   
3:   return  $l_t$ 
```

$$l_t(x) = \log \frac{p(x | z_{1:t})}{1 - p(x | z_{1:t})}$$

*All terms in terms of
"Log-odds ratios"*

Inverse measurement model

*(Typically used when measurements are more complex
than the state)*

- This additive form avoids truncation problems with probabilities close to 1 or 0.



Particle Filters

- “Particles representation” of density:
- Density represented by samples drawn from it
- For a Bayes Filter implementation:

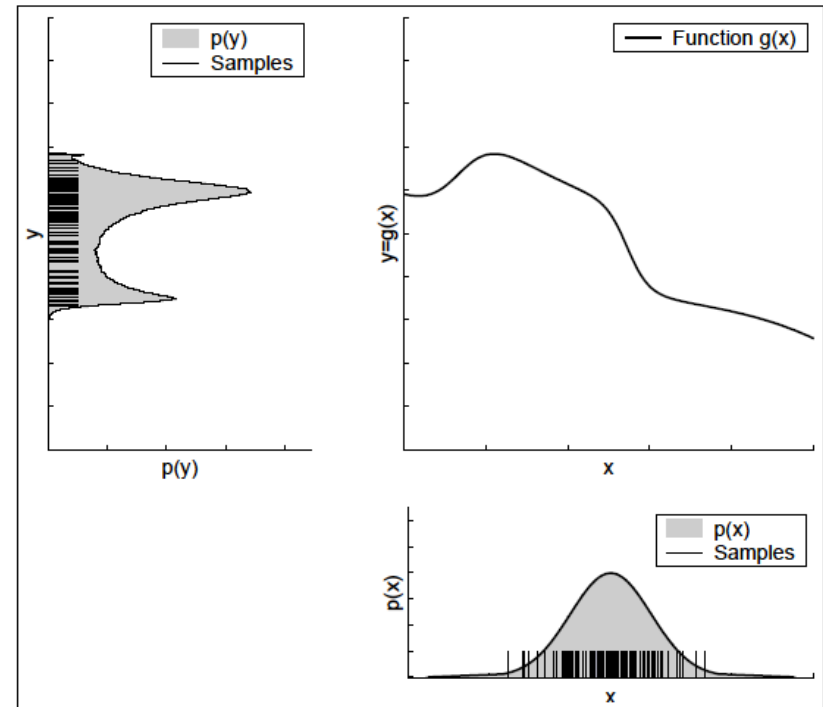
We need to be able to draw samples from:

$p(x'|x, u)$ (motion)

and evaluate:

$p(z|x)$ (sensor)

- Usually easier to do!





Particle Filters

- The “belief” $bel(x_t)$, i.e. the posterior density of robot pose is represented by a particle set:


$$\mathcal{X}_t := x_t^{[1]}, x_t^{[2]}, \dots, x_t^{[M]}$$

- The particle set is recursively updated at each iteration,
- “**Condensation**” of particles around a state indicates “*high posterior likelihood*” of state; given “*measurements*” and “*commands*”



Particle Filters - Example

- How it looks like:





Particle Filters – How it Works?

```
1:   Algorithm Particle_filter( $\mathcal{X}_{t-1}, u_t, z_t$ ):
2:        $\bar{\mathcal{X}}_t = \mathcal{X}_t = \emptyset$ 
3:       for  $m = 1$  to  $M$  do
4:           sample  $x_t^{[m]} \sim p(x_t \mid u_t, x_{t-1}^{[m]})$ 
5:            $w_t^{[m]} = p(z_t \mid x_t^{[m]})$ 
6:            $\bar{\mathcal{X}}_t = \bar{\mathcal{X}}_t + \langle x_t^{[m]}, w_t^{[m]} \rangle$ 
7:       endfor
8:       for  $m = 1$  to  $M$  do
9:           draw  $i$  with probability  $\propto w_t^{[i]}$ 
10:          add  $x_t^{[i]}$  to  $\mathcal{X}_t$ 
11:       endfor
12:       return  $\mathcal{X}_t$ 
```




Particle Filters – How it Works?

```
1: Algorithm Particle_filter( $\mathcal{X}_{t-1}, u_t, z_t$ ):
2:    $\bar{\mathcal{X}}_t = \mathcal{X}_t = \emptyset$ 
3:   for  $m = 1$  to  $M$  do
4:     sample  $x_t^{[m]} \sim p(x_t | u_t, x_{t-1}^{[m]})$ 
5:      $w_t^{[m]} = p(z_t | x_t^{[m]})$ 
6:      $\bar{\mathcal{X}}_t = \bar{\mathcal{X}}_t + \langle x_t^{[m]}, w_t^{[m]} \rangle$ 
7:   endfor
8:   for  $m = 1$  to  $M$  do
9:     draw  $i$  with probability  $w_t^{[m]}$ 
10:    add  $x_t^{[i]}$  to  $\mathcal{X}_t$ 
11:   endfor
12:   return  $\mathcal{X}_t$ 
```

Previous set of “particles”, command and current measurement is used to compute the **intermediate** and **Posterior particle** weights

sets are initialized to empty
We will consider **M** particles

A new sample is drawn from the **motion model**: “prediction step”

A “**weight**” (“**importance factor**”) is assigned for each sample

to integrate the measurement:

measurement model: “Update Step”
New **sample** added to the set, paired with its **weight**



Particle Filters – How it Works?

```
1: Algorithm Particle_filter( $\mathcal{X}_{t-1}, u_t, z_t$ ):
2:    $\bar{\mathcal{X}}_t = \mathcal{X}_t = \emptyset$ 
3:   for  $m = 1$  to  $M$  do
4:     sample  $x_t^{[m]} \sim p(x_t \mid u_t, x_{t-1}^{[m]})$ 
5:      $w_t^{[m]} = p(z_t \mid x_t^{[m]})$ 
6:      $\bar{\mathcal{X}}_t = \bar{\mathcal{X}}_t + \langle x_t^{[m]}, w_t^{[m]} \rangle$ 
7:   endfor
8:   for  $m = 1$  to  $M$  do
9:     draw  $i$  with probability  $\propto w_t^{[i]}$ 
10:    add  $x_t^{[i]}$  to  $\mathcal{X}_t$ 
11:   endfor
12:   return  $\mathcal{X}_t$ 
```

At this point, we have something different than what we started with!

The “real trick” of particle filters:
importance sampling

Samples drawn “with replacement”
i.e., same particles may be picked more than once for the new set
(And some others may be lost)

A new sample set is generated:
Probability of drawing a sample is proportional to its weight



Resampling: Importance Sampling

- **Interesting and necessary step of particle filters,**
- **Transform a set of (x_i, w_i) pairs into a new proper particle set (no weights) for time t ,**
- **Probability of drawing (x_i) should be proportional to its weight (w_i)**
- **M particles are chosen “with replacement”,**
- **(Same particle may be chosen multiple times, some particles may be lost)**



Importance Sampling: Intuition

- When we pass the particles from the motion model, we have the “*prediction*” step,
- We have a set of particles representing:

$$\overline{bel}(x_t)$$

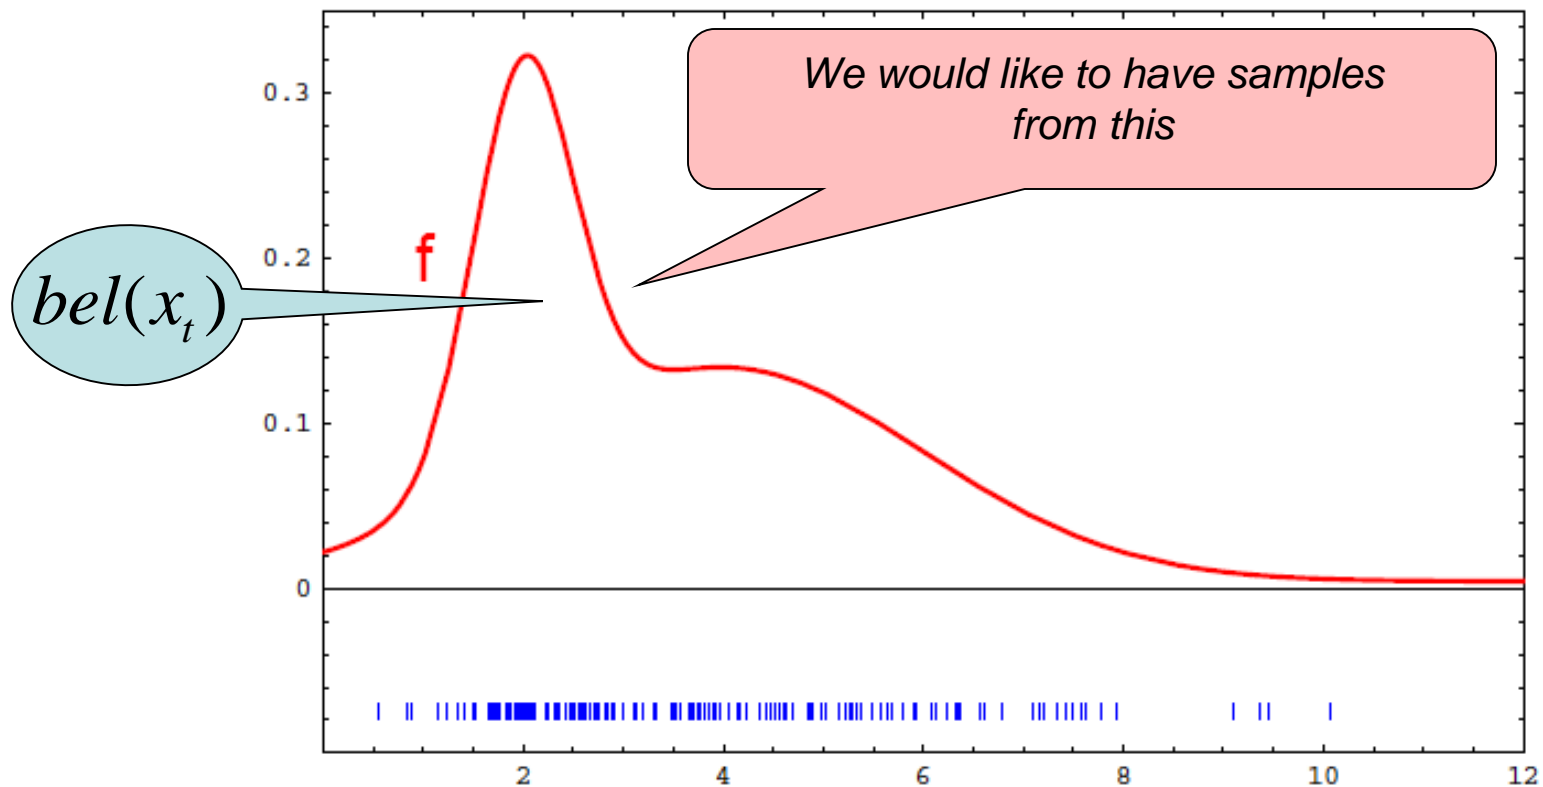
- How can we obtain a set of samples distributed (approximately) according to

$$bel(x_t)$$

which also integrates the measurements?

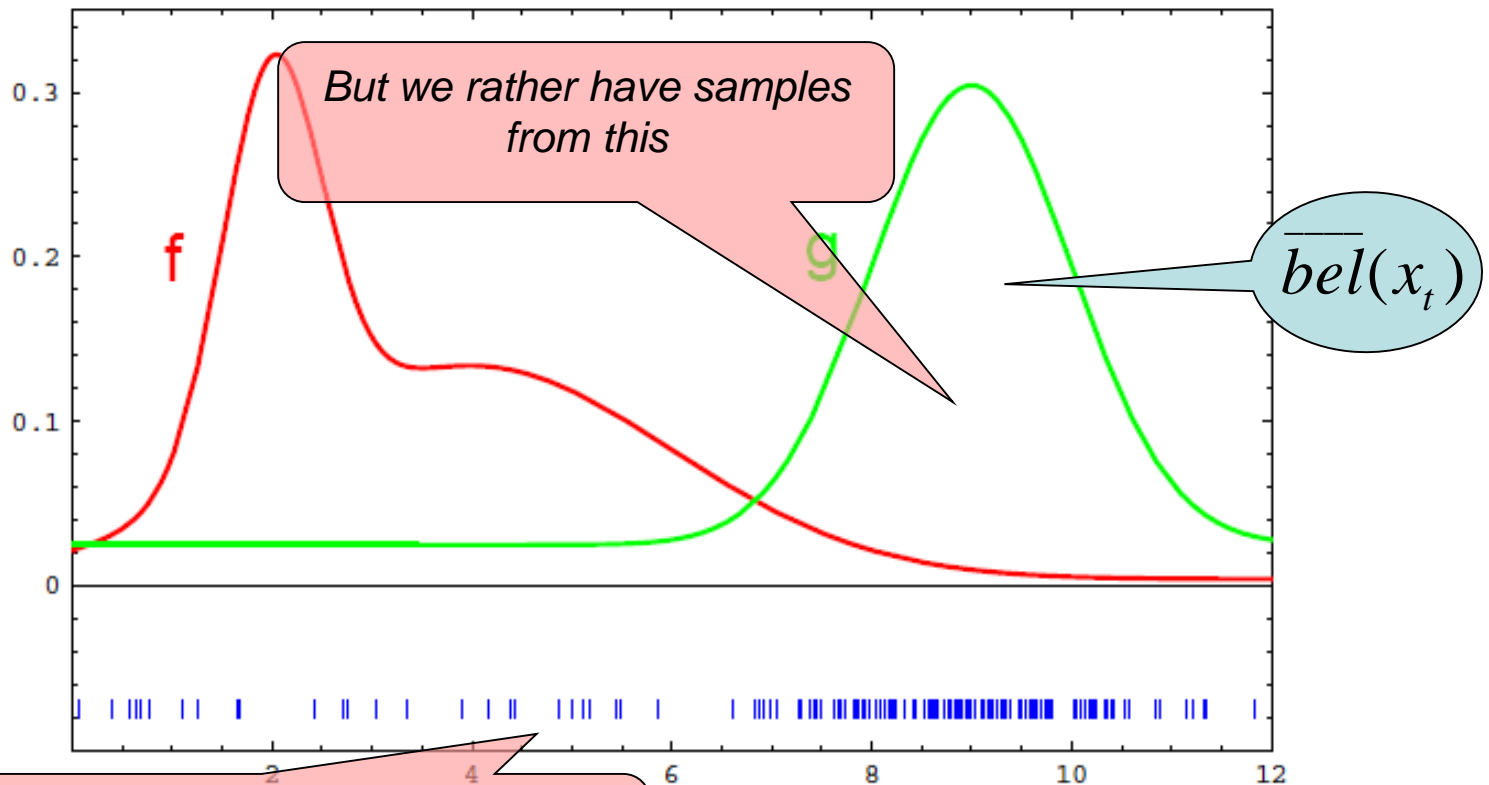


Importance Sampling: Intuition






Importance Sampling: Intuition



Assign weights to these samples according to their "fitness" to the desired density f



Importance Sampling: Intuition



Samples of g can represent f **only** if they are weighted accordingly



Importance Weights

- Consider the update step of the Bayes Filter:

$$bel(x_t) = \eta p(z_t | x_t) \overline{bel}(x_t)$$

- Hence we have:

$$w_i = \frac{f(x_i)}{g(x_i)} = \frac{bel(x_i)}{\overline{bel}(x_i)} = \frac{\text{target distribution}}{\text{proposal distribution}} = \eta p(z_t | x_t)$$


- The normalizer in the w_i (hence its actual value) is irrelevant because we only resample with probabilities *proportional* to w_i .
- Hence we can use in the algorithm:

$$w_i = p(z_t | x_t)$$



The *Resampling Step*

- Now, let us resample according to weights:
- One way: ***Roulette Wheel approach***:



A random number is generated from a uniform distribution.

Probability that a particle is selected is *proportional to its weight*.



Problems With Sampling/Resampling

- **Estimator Variance:**


(Variance of particles as a density estimator)
Is a source of performance loss in Particle Filters

“The fact that statistics (mean, variance, ...) computed from M finite samples drawn from a distribution will differ from the statistics of the original distribution”

That means M samples may be a poor representation of the true density (in particular if M is small).




Estimator Variance Illustration



250 samples

Every particular sample set “instance” results in a different density approximation
(The continuous approximation of particles illustrated with “Kernel density”)



25 samples



Problems With Sampling/Resampling

(Continued)

- **Estimator Variance Divergence**

Major failure mode of original particle filter with Roulette wheel resampling

- **Example:** Stationary robot (constant state) and no sensor (no measurement of state).
- Repetitive resampling monotonically increases estimator variance (particle set converges into a spurious local maxima)
- **End effect:** Particles may be gradually lost (resulting in a single particle) due to random resampling, resulting in “localized” robot (!!!)



Problems With Sampling/Resampling

(Continued)


- **Estimator Variance: One fix**
- **Resampling step may be done less frequently,**
(multiple measurements can still be multiplicatively integrated into the weight factors)
- **Stop resampling if the robot stops / new measurements are not coming,**
- **If resampling is not done, particles may be wasted in regions of low probability,**
- **The proper precision-cost balance requires experience and *tuning* for a particular application.**



Problems With Sampling/Resampling

(Continued)

- **Estimator Variance: Another Fix**
- Use of a “low variance sampler” instead of a roulette wheel strategy
- Use a sequential stochastic process to pick particles



M equally spaced samples.

One random number generates M universally sampled particles

- Stochastic universal sampling
- Systematic resampling
- Linear time complexity
- Easy to implement, low variance



Problems With Sampling/Resampling

(Continued)

- **Particle Deprivation**
- Divergence of the filter characterized by *no samples remaining in the vicinity of the correct state.*
- Result of Estimator Variance problem,
- Can be diminished (but not fixed) by:
 - Increasing sample size (M),
 - Adding uniformly distributed samples at each step,



Again: Particle Filtering Algorithm

```
1:  Algorithm Particle_filter( $\mathcal{X}_{t-1}, u_t, z_t$ ):
2:       $\bar{\mathcal{X}}_t = \mathcal{X}_t = \emptyset$ 
3:      for  $m = 1$  to  $M$  do
4:          sample  $x_t^{[m]} \sim p(x_t \mid u_t, x_{t-1}^{[m]})$ 
5:           $w_t^{[m]} = p(z_t \mid x_t^{[m]})$ 
6:           $\bar{\mathcal{X}}_t = \bar{\mathcal{X}}_t + \langle x_t^{[m]}, w_t^{[m]} \rangle$ 
7:      endfor
8:      for  $m = 1$  to  $M$  do
9:          draw  $i$  with probability  $\propto w_t^{[i]}$ 
10:         add  $x_t^{[i]}$  to  $\mathcal{X}_t$ 
11:      endfor
12:      return  $\mathcal{X}_t$ 
```



Summary

- We discussed non-parametric Bayes Filter implementations: Histogram and Particle Filters,
- Approximation of posterior by finite set of values,
- Histogram Filter: Decompose state space into M *convex regions* – *assign a probability for each region*,
- Decomposition technique plays important role: Static and Dynamic versions, as well as Topological/Non-topo versions determine performance and implementation,
- Particle Filter: Represent the posterior by random samples of state. Recursively update these.
- Very easy implementation and flexible posterior,
- Some failure modes that need special attention